

Master vs Replica Database

1 What is a Master Database? (Purpose)

The **master database** is the **main database** where **all write operations happen**.

Its core responsibility is:

- Handling **INSERT**
- Handling **UPDATE**
- Handling **DELETE**

It is the **source of truth** for the data.

In most systems, **only the master database is allowed to modify data**.

2 What is a Replica Database? (Purpose)

A **replica database** is a **read-only copy** of the master database.

Its main purpose is:

- Handling **read operations (SELECT)**
- Reducing load on the master
- Improving read performance

The replica keeps syncing data from the master using **replication mechanisms**.

3 Why Do We Need Master–Replica Setup?

In real software systems:

- Writes are **less frequent**
- Reads are **very frequent**

If everything hits one database:

- Performance degrades
- Database becomes a bottleneck

So we separate responsibilities:

- Master → writes
- Replica → reads

What Happens Behind the Scenes? (Important)

When data is written:

1. Master updates its data
2. Changes are sent to replica
3. Replica applies changes

⚠ Replication is usually **not instant** (called **replication lag**).

That means:

- A read immediately after a write **might not reflect new data**

Connection Pool:

What you got right

Yes:

- A **connection pool** is a set of database connections
- These connections are **created and managed once**
- When a request needs DB access → it **borrowes a connection**
- When done → the connection is **returned to the pool**
- Connections are **reused**, not destroyed

So this mental model is correct 👍

Request → get connection → run query → return connection

✖ Small Correction (Very Important)

“**No other connections can be built except of pools**”

This part is **not strictly true**.

Correct version:

- **Connection pool limits how many connections SHOULD exist**
- But technically:
 - You *can* create connections outside the pool

- You *should not* in a well-designed backend system
-

Correct & Professional Statement (Use This)

A connection pool manages a limited number of database connections and reuses them for incoming requests. New connections are created only by the pool, and requests must wait if all connections are in use.

What Actually Happens in Practice

Pool Rules:

- Pool has a **maximum size** (e.g. 10 connections)
- Pool creates connections **lazily or at startup**
- Requests must **borrow** from the pool

If all connections are busy:

- Request **waits**
- OR times out
- OR fails (depending on config)

No unlimited connection creation.

Why This Rule Exists (System Design Reason)

Databases have:

- Limited connection capacity
- Memory and CPU cost per connection

So pooling:

- Protects the database
- Controls concurrency
- Prevents crashes

SQLAlchemy Engine (Built on Connection Pool Concept)

1 What is SQLAlchemy Engine? (Main Purpose)

The SQLAlchemy Engine is the **central object** that:

- Manages **database connections**
- Uses a **connection pool internally**
- Executes SQL queries
- Acts as the **bridge between Python code and the database**

👉 You never talk to the database directly

👉 You talk to the Engine

2 How Engine Relates to Connection Pool (Key Idea)

When you create an engine:

```
engine = create_engine("postgresql://user:pass@localhost/db")
```

What actually happens:

- SQLAlchemy creates a **connection pool**
- Connections are created **lazily**
- Engine becomes the **manager of that pool**

So:

Engine = Pool + Dialect + DBAPI

3 What the Engine Does Internally

Behind the scenes:

- Keeps a **pool of open connections**
- Gives a connection when needed
- Takes it back when done
- Prevents unlimited connections

You **never see the pool**, but it's always there.

What happens internally:

1. engine.connect() → asks pool for a connection
2. Query runs
3. with block ends → connection goes back to pool

No manual open/close.

Engine with Explicit Pool Configuration

```
engine = create_engine(  
    "postgresql://user:pass@localhost/db",  
    pool_size=5,  
    max_overflow=10,  
    pool_timeout=30  
)
```

Meaning:

- pool_size=5 → always keep 5 connections
- max_overflow=10 → allow up to 10 extra temporarily
- pool_timeout=30 → wait 30 seconds if pool is full

This is **exactly your pool concept**, now formalized.

Session Factory (SQLAlchemy)

1 What is a Session Factory? (Main Purpose)

A **session factory** is an object that **creates Session instances** when needed.

It does **not** talk to the database itself.

Its only job is:

👉 “Whenever the application needs to interact with the database, create a new Session configured with the Engine.”

Think of it as a **Session generator**, not a Session.

2 Why Session Factory Exists (Very Important)

In a backend application:

- Multiple requests happen at the same time
- Each request must have its **own session**
- Sessions **must not be shared** across requests

So instead of manually creating sessions:

- We define **one factory**
 - Factory creates **independent sessions**
-

3 How Session Factory is Created (Python)

```
from sqlalchemy.orm import sessionmaker  
  
SessionLocal = sessionmaker(bind=engine)
```

Here:

- SessionLocal is the **session factory**
 - engine provides access to the connection pool
-

4 How Session Factory is Used

```
session = SessionLocal()  
  
users = session.query(User).all()  
  
session.close()
```

Flow:

1. Factory creates a new session
 2. Session borrows a connection from Engine
 3. Queries execute
 4. Session closes → connection returned to pool
-

5 How Session Factory Fits in the Big Picture

```
Request
  ↓
Session (created by factory)
  ↓
Engine
  ↓
Connection Pool
  ↓
Database
```

The **factory sits above the Engine**, not below it.

6 Session Factory vs Session (Critical Difference)

Session Factory	Session
Creates sessions	Executes queries
One per application	One per request
Stateless	Stateful
Long-lived	Short-lived

Clean Explanation (Professional Version)

The SQLAlchemy Engine is created once at application startup.
When the engine is created, it also initializes a connection pool.
This pool persists for the lifetime of the application.

A session factory is then bound to this engine.

For every request:

- A new Session is created by the session factory
- The Session uses the Engine
- The Engine provides a connection from the existing pool
- After the request finishes, the Session is closed
- The connection is returned to the pool, not destroyed

So:

Engine & Pool = long-lived

Session = short-lived (per request)

Visual Flow (Board-Friendly)

App Startup:
Engine created
↓
Connection Pool created

Runtime:
Request 1 → Session → Engine → Pool → DB → return connection
Request 2 → Session → Engine → Pool → DB → return connection
Request 3 → Session → Engine → Pool → DB → return connection

Key Sentence (Simple & Accurate)

The engine creates and owns the connection pool once, and all sessions reuse that same pool through the engine.

Even Simpler Version (Beginner Language)

The engine is created once.
The connection pool is created once.
Every request only creates a new session, not new connections.
Sessions reuse the same pool through the engine.

Python Structure (Concept Mapping)

```
# created once
engine = create_engine(DB_URL)

# created once
SessionLocal = sessionmaker(bind=engine)

# per request
session = SessionLocal()
```

Only **session** is per request.
Engine and pool are **global/shared**.

Interview-Grade One-Liner

The engine initializes the connection pool once at startup, and all sessions created per request reuse that pool via the engine.

If Someone Tries to Trap You

Q: “Does every session create a new connection pool?”

Answer:

No. The pool is created by the engine once. Sessions only borrow connections from the existing pool.

Multi-Tenancy

1 What is Multi-Tenancy? (Main Purpose)

Multi-tenancy is an architecture where **one application instance serves multiple independent customers (tenants)**, while **keeping each tenant's data isolated**.

A *tenant* can be:

- A company
- An organization
- A team
- A client account

Main purpose:

👉 Serve multiple tenants using shared infrastructure while isolating data.

2 Why Multi-Tenancy Exists

Without multi-tenancy:

- You would deploy **one application per customer**
- Separate databases, servers, configs
- High cost, hard maintenance

With multi-tenancy:

- One backend
 - One codebase
 - Controlled data separation
-

3 Core Rule of Multi-Tenancy (Must Remember)


Tenants share the application but not each other's data.

Isolation is **logical**, not accidental.

4 Common Multi-Tenancy Models (Software-Only)

◆ 1. Database-per-Tenant

Each tenant has its **own database**.




Use when:

- Strong isolation required
 - Enterprise clients
 - Custom scaling per tenant
-

◆ 2. Schema-per-Tenant

One database, multiple schemas.



Use when:

- Medium isolation
 - Controlled environment
 - PostgreSQL systems
-

◆ 3. Table-level (Shared Tables)

All tenants share tables, data separated by `tenant_id`.

```
users
-----
id | name | tenant_id
```

Use when:

- Large number of tenants
- SaaS products

- Cost efficiency matters
-

5 Python Example (Table-Level Multi-Tenancy)

```
def get_users(session, tenant_id):  
    return session.query(User).filter(  
        User.tenant_id == tenant_id  
    ).all()
```

Here:

- Same table
 - Same database
 - Data separated by `tenant_id`
-

6 How Multi-Tenancy Works with SQLAlchemy (Important)

At request start:

1. Identify tenant (from token, subdomain, header)
2. Configure session/query to use tenant context
3. Ensure **every query is tenant-scoped**

Example:

```
tenant_id = request.tenant_id  
session.query(Order).filter(Order.tenant_id == tenant_id)
```

7 Multi-Tenancy with Multiple Databases (Advanced)

```
engine = create_engine(DB_URL_TENANT_A)  
SessionLocal = sessionmaker(bind=engine)
```

Tenant decides:

- Which engine
- Which database

Engine & pool still follow the same rules you learned earlier.

8 Multi-Tenancy + Connection Pool (Key Insight)

- Engine created per database
 - Each engine has its **own pool**
 - Sessions reuse pools
 - Tenant routing happens **before session creation**
-

9 Use Cases in Software Engineering

- SaaS platforms
- CRM systems
- HR management systems
- Analytics dashboards
- ERP systems
- Multi-company admin panels

Basically:

👉 Any product serving multiple organizations

10 Advantages

- ✓ Cost efficient
 - ✓ Single deployment
 - ✓ Easy updates
 - ✓ Scalable architecture
-

1 1 Challenges (Must Mention)

- ✗ Data isolation bugs
- ✗ Security risks if tenant filtering fails
- ✗ Complex queries
- ✗ Migration complexity

1 2 Interview-Ready Definition

Multi-tenancy is an architectural approach where a single application serves multiple tenants while keeping their data logically isolated.

What this code is (one line)

This is an async database connection manager that sets up master–replica PostgreSQL engines, connection pools, and separate session factories for read and write operations.

High-Level What It Does

1. Creates **one async SQLAlchemy engine for master**
2. Optionally creates **one async engine for replica**
3. Each engine **creates its own connection pool once**
4. Creates **two session factories**:
 - o Write sessions → master
 - o Read sessions → replica (fallback to master)
5. Provides a **clean API** to get the correct session per request
6. Handles **startup, shutdown, and health checks**