

Step-by-Step (Mapped to Concepts You Learned)

1 Engine & Connection Pool (Created Once)

```
self.master_engine = create_async_engine(...)  
self.replica_engine = create_async_engine(...)
```

- Each `create_async_engine`:
 - Creates an **engine**
 - Initializes a **connection pool**
- Pools live for the **lifetime of the application**

✓ This matches: *engine creates pool once*

2 Master–Replica Separation

```
self.master_engine    → writes  
self.replica_engine  → reads (if available)
```

- Writes **always go to master**
- Reads **prefer replica**
- If replica fails → fallback to master

✓ Correct master–replica design

3 Session Factories (Per-Request Sessions)

```
self.write_session_factory = async_sessionmaker(bind=self.master_engine)  
self.read_session_factory = async_sessionmaker(bind=read_engine)
```

- Factories are created **once**
- Every request gets a **new session**
- Sessions borrow connections from the engine's pool

✓ Correct session lifecycle

4 Read / Write Routing Logic

```
def get_session(operation_type="write"):
```

```
if operation_type == "read":  
    return self.read_session_factory()  
return self.write_session_factory()
```

- Read request → replica session
- Write request → master session

✓ Simple and safe routing

5 Connection Testing at Startup

```
SELECT 1
```

- Ensures:
 - Master is reachable
 - Replica is reachable
- If replica fails → disabled automatically

✓ Prevents runtime crashes

6 Graceful Shutdown

```
await engine.dispose()
```

- Closes all pooled connections
- Clean app shutdown

✓ Correct async cleanup

How a Request Uses This (Flow)

```
Request  
↓  
get_session("read" / "write")  
↓  
AsyncSession (new)  
↓  
Engine (shared)  
↓  
Connection Pool (shared)  
↓  
Database
```

Only the **session** is per request.
Engines & pools are **shared and reused**

Database Replication (Master → Replica)

Replication is the **process of copying changes from the master database to one or more replicas** automatically so all databases stay in sync.

1 Types of Replication (Software Context)

1. Synchronous Replication

- Master **waits for replica to confirm** that data is written before committing.
- Guarantees **strong consistency**.
- Slower, because writes wait for replica acknowledgement.
- Use case: **financial apps, critical systems**.

2. Asynchronous Replication

- Master **does not wait** for replica.
- Replica gets updates **with a delay** (replication lag).
- Faster, but reads immediately after a write might not see latest data.
- Use case: **high-traffic apps, dashboards, SaaS platforms**.

2 How Replication Works Internally

Think **log-based replication**, which is most common in PostgreSQL / MySQL:

1. **Master keeps a write-ahead log (WAL)**
 - Every INSERT, UPDATE, DELETE is recorded in a log.
2. **Replica reads the WAL**
 - Applies the same changes in order.
3. **Replica applies changes to its database**
 - Keeps schema and data consistent with master.

Python backend doesn't handle replication itself — it just talks to master/replica engines.

3 Replication Flow (Step-by-Step)

1. Backend writes → Master DB
2. Master writes → WAL (write-ahead log)
3. Replica reads WAL (continuously)
4. Replica applies changes
5. Reads on replica now reflect master's data (after lag)

1 How it actually works:

- **Every session is created per request** (or per operation).
- **You pass a tenant_id when creating the session.**
- That session “**carries**” that tenant_id so all queries in that session know which tenant’s data to access.

So yes: **every session is tied to one tenant_id.**

The CORRECT way to explain it (say this)

“I achieved multi-tenancy by creating **separate database engines per tenant**.

Each tenant has its **own connection pool** and **its own database**.

When a request comes in, the tenant ID is used to **select the correct engine**, and a session is created from that tenant’s connection pool.”

That statement is **100% accurate**.

Slightly more detailed (if they ask “how exactly?”)

“At application startup or when a tenant is registered, I create a dedicated master and replica engine for that tenant.

These engines maintain their own connection pools.

During runtime, every database request includes a tenant ID, and the database manager uses that tenant ID to route the request to the correct engine and database.”

Important correction (small but critical)

You should **not** say:

“the session provides him access to the pool”

Say this instead:

“**the session is created using the tenant’s engine, which internally manages its own connection pool.**”

Why?

- Session does not *own* the pool
- Engine owns the pool
- Session borrows a connection from the pool

That distinction sounds **very professional**.

One-sentence ultra-polished version (INTERVIEW GOLD)

Memorize this 

“Multi-tenancy is achieved by dynamically routing database sessions based on tenant ID, where each tenant has its own database engine and connection pool, ensuring complete data isolation.”

If they ask “what ensures isolation?”

Answer:

“Isolation is guaranteed because each tenant connects to a separate database using a separate engine and pool, so there is no shared storage or shared connections between tenants.”

If they ask “what is shared then?”

Answer:

“Only the application layer is shared — the codebase, configuration logic, and database manager — not the data.”

Final confidence check

Your understanding is **correct**:

- tenant → own engine 
- engine → own pool 
- tenant ID → routing key 
- session → created from tenant engine 
- database → isolated 

1. **What to say first (simple)**
2. **How to explain it using your current implementation**

-
- 3. How to explain it in a real-world / production way
 - 4. One-line interview version
-

1 Short, simple answer (say this first)

“Tenants are registered by creating a dedicated database for the tenant and then initializing database engines for that tenant inside the application.”

That’s enough to start.

2 How tenants are registered in your current implementation

Be honest and clear (this is important):

“In my current implementation, tenant registration is done during application setup. I manually define tenant IDs and database names, create the databases, and then call a method that connects the tenant to the application.”

Then explain **what that method does**:

“This method creates a master and replica engine for the tenant and stores them in a tenant-engine registry inside the database manager.”

This shows you understand the **lifecycle**.

3 How tenant registration would work in a real system (VERY IMPORTANT)

This is where you level up.

Say this:

“In a real production system, tenant registration would be automated.”

Then explain step-by-step in easy words:

“When a new tenant signs up:

1. A tenant record is created in a central metadata store
2. A new database is created for that tenant
3. Database migrations or table creation are applied
4. A database engine and connection pool are initialized
5. The tenant is added to the tenant registry
6. From that point onward, requests containing that tenant ID are routed to that database”

This maps perfectly to your current design.

4 Where tenant information is stored

If they ask “*where do you store tenants?*” say:

“Tenant metadata is stored separately from tenant data.

The application only uses the tenant ID to resolve the correct database connection.”

You can also add:

“This avoids mixing tenant configuration with tenant business data.”

5 If they ask “how does the request know the tenant?”

Say this:

“Each request carries a tenant identifier, usually through headers, subdomain, or authentication context.

The backend extracts the tenant ID and uses it to select the correct database connection.”

Even if you didn’t implement headers yet — this is the **standard explanation**.

6 One-line interview-ready answer (MEMORIZE THIS)

“Tenants are registered by provisioning a dedicated database for each tenant and initializing corresponding database engines, after which the tenant ID is used at runtime to route requests to the correct database.”

7 If they challenge you: “Is tenant registration dynamic?”

Answer confidently:

“The design supports dynamic tenant registration.

In my demo it’s static for simplicity, but the same mechanism can be triggered at runtime without changes to the core architecture.”

 That answer shows **foresight**.